

Dimensionality Reduction

1. Problem Definition.

Carrefour Kenya is a part of the French multination Carrefour S.A. The company specialises in retail, operating in several shopping malls in the country including The Hub, Two Rivers and THika Road mall among others. The company also operates using several e-retail platforms mainly the company's website as well as through Jumia market place.

The retail seeks to roll out several marketing strategies to increase their sales numbers. The marketing team is seeking the help of the company's data scientist to assist in determining which are the best strategies to employ that will result in the highest number of sales.

2. Data Sourcing

The data used for the analysis was sourced from here. The data contains a list of transaction from various customers, it includes the type of customer, whether a Normal customer or a Member, their gender, product line they purchased, quantities purchased, etc.

3. Loading and Checking the Data

```
carre = read.csv("http://bit.ly/CarreFourDataset")
```

```
head(carre)
```

##	Invoice.ID	Branch	Customer.type	Gender	Product.line	Unit.price	
## 1	750-67-8428	A	Member	Female	Health and beauty	74.69	
## 2	226-31-3081	C	Normal	Female	Electronic accessories	15.28	
## 3	631-41-3108	A	Normal	Male	Home and lifestyle	46.33	
## 4	123-19-1176	A	Member	Male	Health and beauty	58.22	
## 5	373-73-7910	A	Normal	Male	Sports and travel	86.31	
## 6	699-14-3026	C	Normal	Male	Electronic accessories	85.39	
##	Quantity	Tax	Date	Time	Payment	cogs	gross.margin.percentage
## 1	7	26.1415	1/5/2019	13:08	Ewallet	522.83	4.761905
## 2	5	3.8200	3/8/2019	10:29	Cash	76.40	4.761905
## 3	7	16.2155	3/3/2019	13:23	Credit card	324.31	4.761905
## 4	8	23.2880	1/27/2019	20:33	Ewallet	465.76	4.761905
## 5	7	30.2085	2/8/2019	10:37	Ewallet	604.17	4.761905
## 6	7	29.8865	3/25/2019	18:30	Ewallet	597.73	4.761905
##	gross.income	Rating	Total				
## 1	26.1415	9.1	548.9715				
## 2	3.8200	9.6	80.2200				
## 3	16.2155	7.4	340.5255				
## 4	23.2880	8.4	489.0480				
## 5	30.2085	5.3	634.3785				
## 6	29.8865	4.1	627.6165				

```
tail(carre)
```

##	Invoice.ID	Branch	Customer.type	Gender	Product.line	Unit.price
----	------------	--------	---------------	--------	--------------	------------

```
## 995 652-49-6720      C      Member Female Electronic accessories      60.95
## 996 233-67-5758      C      Normal   Male   Health and beauty      40.35
## 997 303-96-2227      B      Normal Female   Home and lifestyle      97.38
## 998 727-02-1313      A      Member   Male   Food and beverages      31.84
## 999 347-56-2442      A      Normal   Male   Home and lifestyle      65.82
## 1000 849-09-3807     A      Member Female   Fashion accessories      88.34
##      Quantity      Tax      Date Time Payment      cogs gross.margin.percentage
## 995      1  3.0475 2/18/2019 11:40 Ewallet  60.95      4.761905
## 996      1  2.0175 1/29/2019 13:46 Ewallet  40.35      4.761905
## 997     10 48.6900 3/2/2019 17:16 Ewallet 973.80      4.761905
## 998      1  1.5920 2/9/2019 13:22   Cash  31.84      4.761905
## 999      1  3.2910 2/22/2019 15:33   Cash  65.82      4.761905
## 1000     7 30.9190 2/18/2019 13:28   Cash 618.38      4.761905
##      gross.income Rating      Total
## 995      3.0475      5.9    63.9975
## 996      2.0175      6.2    42.3675
## 997     48.6900      4.4  1022.4900
## 998      1.5920      7.7    33.4320
## 999      3.2910      4.1    69.1110
## 1000     30.9190      6.6   649.2990
```

```
dim(carre)
```

```
## [1] 1000  16
```

Our dataset is made up of 16 columns and 1000 rows

```
# checking the data type as we will need to change some when performing dimensionality reduction
apply(carre, class)
```

```
##      Invoice.ID      Branch      Customer.type
##      "factor"      "factor"      "factor"
##      Gender      Product.line      Unit.price
##      "factor"      "factor"      "numeric"
##      Quantity      Tax      Date
##      "integer"      "numeric"      "factor"
##      Time      Payment      cogs
##      "factor"      "factor"      "numeric"
## gross.margin.percentage      gross.income      Rating
##      "numeric"      "numeric"      "numeric"
##      Total
##      "numeric"
```

4. Data Cleaning

Checking for and Dealing with Null Values

```
colSums(is.na(carre))
```

```
##      Invoice.ID      Branch      Customer.type
##      0      0      0
```

```
##           Gender           Product.line           Unit.price
##           0              0              0
##           Quantity          Tax              Date
##           0              0              0
##           Time           Payment           cogs
##           0              0              0
## gross.margin.percentage    gross.income           Rating
##           0              0              0
##           Total
##           0
```

Our dataset contains no null values.

Checking for and Dealing with Duplicates

```
duplicates = carre[duplicated(carre),]
dim(duplicates)
```

```
## [1] 0 16
```

Our dataset contains no duplicate values.

5. Dimensionality Reduction

PCA

Data Preparation

```
library(CatEncoders)
```

```
## Warning: package 'CatEncoders' was built under R version 3.6.3
```

```
##
## Attaching package: 'CatEncoders'
```

```
## The following object is masked from 'package:base':
##
##      transform
```

```
#Changing categorical data to numerical data
factors <- names(which(sapply(carre, is.factor)))
# Label Encoder
for (i in factors){
  encode <- LabelEncoder.fit(carre[, i])
  carre[, i] <- transform(encode, carre[, i])
}
```

```
#checking whether all columns have changed to the correct data types
```

```
sapply(carre, class)
```

```
## Invoice.ID Branch Customer.type
## "integer" "integer" "integer"
## Gender Product.line Unit.price
## "integer" "integer" "numeric"
## Quantity Tax Date
## "integer" "numeric" "integer"
## Time Payment cogs
## "integer" "integer" "numeric"
## gross.margin.percentage gross.income Rating
## "numeric" "numeric" "numeric"
## Total
## "numeric"
```

```
#dropping the invoice ID column
```

```
carre$Invoice.ID <- NULL
```

```
#chaecking our new dataset
```

```
head(carre)
```

```
## Branch Customer.type Gender Product.line Unit.price Quantity Tax Date
## 1 1 1 1 4 74.69 7 26.1415 27
## 2 3 2 1 1 15.28 5 3.8200 88
## 3 1 2 2 5 46.33 7 16.2155 82
## 4 1 1 2 4 58.22 8 23.2880 20
## 5 1 2 2 6 86.31 7 30.2085 58
## 6 3 2 2 1 85.39 7 29.8865 77
## Time Payment cogs gross.margin.percentage gross.income Rating Total
## 1 147 3 522.83 4.761905 26.1415 9.1 548.9715
## 2 24 1 76.40 4.761905 3.8200 9.6 80.2200
## 3 156 2 324.31 4.761905 16.2155 7.4 340.5255
## 4 486 3 465.76 4.761905 23.2880 8.4 489.0480
## 5 30 3 604.17 4.761905 30.2085 5.3 634.3785
## 6 394 3 597.73 4.761905 29.8865 4.1 627.6165
```

```
# carrying out our PCA to check for
```

```
carre.pca <- prcomp(carre, center = TRUE, scale. = FALSE)
```

```
# previewing our PCA summary
```

```
summary(carre.pca)
```

```
## Importance of components:
```

```
## PC1 PC2 PC3 PC4 PC5 PC6
## Standard deviation 340.384 147.0671 25.88456 20.51561 1.73430 1.69601
## Proportion of Variance 0.836 0.1561 0.00483 0.00304 0.00002 0.00002
## Cumulative Proportion 0.836 0.9921 0.99690 0.99993 0.99995 0.99998
## PC7 PC8 PC9 PC10 PC11 PC12 PC13
## Standard deviation 1.24606 0.84341 0.7999 0.5068 0.487 7.832e-14 1.527e-14
## Proportion of Variance 0.00001 0.00001 0.0000 0.0000 0.000 0.000e+00 0.000e+00
## Cumulative Proportion 0.99999 0.99999 1.0000 1.0000 1.000 1.000e+00 1.000e+00
## PC14 PC15
## Standard deviation 8.532e-16 6.292e-31
```

```
## Proportion of Variance 0.000e+00 0.000e+00
## Cumulative Proportion 1.000e+00 1.000e+00
```

```
str(carre.pca)
```

```
## List of 5
## $ sdev      : num [1:15] 340.38 147.07 25.88 20.52 1.73 ...
## $ rotation: num [1:15, 1:15] 9.85e-05 -2.89e-05 -7.26e-05 1.60e-04 4.95e-02 ...
##   .- attr(*, "dimnames")=List of 2
##   .. ..$ : chr [1:15] "Branch" "Customer.type" "Gender" "Product.line" ...
##   .. ..$ : chr [1:15] "PC1" "PC2" "PC3" "PC4" ...
## $ center    : Named num [1:15] 1.99 1.5 1.5 3.45 55.67 ...
##   .- attr(*, "names")= chr [1:15] "Branch" "Customer.type" "Gender" "Product.line" ...
## $ scale      : logi FALSE
## $ x          : num [1:1000, 1:15] 313 -337 24 229 432 ...
##   .- attr(*, "dimnames")=List of 2
##   .. ..$ : NULL
##   .. ..$ : chr [1:15] "PC1" "PC2" "PC3" "PC4" ...
## - attr(*, "class")= chr "prcomp"
```


```
# Installing our ggbiplot visualisation package
#
library(devtools)
```

```
## Loading required package: usethis
```

```
#visualising our PCA
```

```
plot(carre.pca, type="l")
```

carre.pca



t-Distributed Stochastic Neighbor Embedding (t-SNE)

```
library(Rtsne)
```

```
## Warning: package 'Rtsne' was built under R version 3.6.3
```

```
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 3.6.3
```

```
## corrplot 0.84 loaded
```

```
library(lattice)
```

Trying t-SNE with Branch


```
carre.branch <- carre  
carre.branch$Branch <- as.factor(carre.branch$Branch)
```

```
colors = rainbow(length(unique(carre.branch$Branch)))
names(colors) = unique(carre.branch$Branch)
```

```
tsne <- Rtsne(carre[, -1], dims = 2, perplexity=30, verbose=TRUE, max_iter = 500)
```

```
## Performing PCA
## Read the 1000 x 14 data matrix successfully!
## OpenMP is working. 1 threads.
## Using no_dims = 2, perplexity = 30.000000, and theta = 0.500000
## Computing input similarities...
## Building tree...
## Done in 0.19 seconds (sparsity = 0.103660)!
## Learning embedding...
## Iteration 50: error is 64.385102 (50 iterations in 0.21 seconds)
## Iteration 100: error is 57.876485 (50 iterations in 0.13 seconds)
## Iteration 150: error is 57.615991 (50 iterations in 0.13 seconds)
## Iteration 200: error is 57.590925 (50 iterations in 0.12 seconds)
## Iteration 250: error is 57.586959 (50 iterations in 0.12 seconds)
## Iteration 300: error is 0.752410 (50 iterations in 0.11 seconds)
## Iteration 350: error is 0.609841 (50 iterations in 0.12 seconds)
## Iteration 400: error is 0.579446 (50 iterations in 0.12 seconds)
## Iteration 450: error is 0.566818 (50 iterations in 0.11 seconds)
## Iteration 500: error is 0.558625 (50 iterations in 0.13 seconds)
## Fitting performed in 1.30 seconds.
```

```
plot(tsne$Y, t='n', main="tsne")
text(tsne$Y, labels=carre$Branch, col=colors[carre$Branch])
```



Trying T-SNE with Product Line

```
carre.prod <- carre
carre.prod$Product.line <- as.factor(carre.prod$Product.line)

# for plotting
colors = rainbow(length(unique(carre.prod$Product.line)))
names(colors) = unique(carre.prod$Product.line)

tsne <- Rtsne(carre[, -4], dims = 2, perplexity = 30, verbose = TRUE,
              max_iter = 500)
```

```
## Performing PCA
## Read the 1000 x 14 data matrix successfully!
## OpenMP is working. 1 threads.
## Using no_dims = 2, perplexity = 30.000000, and theta = 0.500000
## Computing input similarities...
## Building tree...
## Done in 0.17 seconds (sparsity = 0.103656)!
## Learning embedding...
## Iteration 50: error is 63.981024 (50 iterations in 0.17 seconds)
## Iteration 100: error is 57.979898 (50 iterations in 0.17 seconds)
## Iteration 150: error is 57.710602 (50 iterations in 0.16 seconds)
## Iteration 200: error is 57.681837 (50 iterations in 0.15 seconds)
## Iteration 250: error is 57.677603 (50 iterations in 0.17 seconds)
## Iteration 300: error is 0.753584 (50 iterations in 0.13 seconds)
## Iteration 350: error is 0.610855 (50 iterations in 0.12 seconds)
```



```
## Iteration 400: error is 0.575060 (50 iterations in 0.13 seconds)
## Iteration 450: error is 0.562386 (50 iterations in 0.13 seconds)
## Iteration 500: error is 0.555600 (50 iterations in 0.13 seconds)
## Fitting performed in 1.45 seconds.
```

```
# plotting our graph
plot(tsne$Y, t = 'n', main = 'tsne')
text(tsne$Y, labels = carre.prod$Product.line,
col = colors[carre.prod$Product.line])
```

